


```
In [1]: ▶ import pandas as pd
import numpy as np
import seaborn as sns
import plotly
import matplotlib.pyplot as plt
import sklearn
from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
```

```
In [2]: ▶ df = pd.read_csv('C:/Users/User/Desktop/datasciencce/employee_promotion.csv')
```

```
In [3]: ▶ df.head()
```

Out[3]:

	employee_id	department	region	education	gender	recruitment_channel	no_of_training
0	65438	Sales & Marketing	region_7	Master's & above	f	sourcing	
1	65141	Operations	region_22	Bachelor's	m	other	
2	7513	Sales & Marketing	region_19	Bachelor's	m	sourcing	
3	2542	Sales & Marketing	region_23	Bachelor's	m	other	
4	48945	Technology	region_26	Bachelor's	m	other	



```
In [4]: ▶ df.shape
```

Out[4]: (54808, 13)


```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54808 entries, 0 to 54807
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   employee_id           54808 non-null  int64  
 1   department            54808 non-null  object  
 2   region               54808 non-null  object  
 3   education             52399 non-null  object  
 4   gender               54808 non-null  object  
 5   recruitment_channel    54808 non-null  object  
 6   no_of_trainings       54808 non-null  int64  
 7   age                  54808 non-null  int64  
 8   previous_year_rating  50684 non-null  float64 
 9   length_of_service     54808 non-null  int64  
10   awards_won           54808 non-null  int64  
11   avg_training_score    52248 non-null  float64 
12   is_promoted          54808 non-null  int64  
dtypes: float64(2), int64(6), object(5)
memory usage: 5.4+ MB
```


```
In [6]: df['department'].value_counts()
```

```
Out[6]: Sales & Marketing    16840
Operations    11348
Technology    7138
Procurement  7138
Analytics    5352
Finance      2536
HR           2418
Legal        1039
R&D           999
Name: department, dtype: int64
```


```
In [7]: plt.figure(figsize=(15,5))
df['department'].value_counts().plot(kind='barh',color='purple')
plt.show()
```




```
In [8]: ▶ plt.figure(figsize=(10,5))
df['education'].value_counts().plot(kind='barh',color='orange')
plt.show()
```



```
In [9]: ▶ plt.figure(figsize=(10,5))
df['gender'].value_counts().plot(kind='bar',color='green')
plt.xticks(rotation=0)
plt.show()
```



```
In [10]: ▶ plt.figure(figsize=(10,5))
df['recruitment_channel'].value_counts().plot(kind='bar',color='blue')
plt.xticks(rotation=0)
plt.show()
```



Encoding

```
In [11]: ▶ le = LabelEncoder()
```

```
In [12]: ▶ df['department']=le.fit_transform(df['department'])
```

```
In [13]: ▶ df['region']=le.fit_transform(df['region'])
df['education']=le.fit_transform(df['education'])
df['recruitment_channel']=le.fit_transform(df['recruitment_channel'])
df['gender']=le.fit_transform(df['gender'])
```

In [14]: `df.head()`

Out[14]:

	employee_id	department	region	education	gender	recruitment_channel	no_of_trainings
0	65438	7	31	2	0	2	1
1	65141	4	14	0	1	0	1
2	7513	7	10	0	1	2	1
3	2542	7	15	0	1	0	2
4	48945	8	18	0	1	0	1



In [15]: `df = df.drop('employee_id',axis=1)`

In [16]: `cols = ['previous_year_rating', 'avg_training_score']`
`for col in cols:`
`df[col] = df[col].apply(lambda x: int(x) if x == x else 0)`

In [17]: `df.head()`

Out[17]:

	department	region	education	gender	recruitment_channel	no_of_trainings	age	previous
0	7	31	2	0	2	1	35	
1	4	14	0	1	0	1	30	
2	7	10	0	1	2	1	34	
3	7	15	0	1	0	2	39	
4	8	18	0	1	0	1	45	



In [18]: `df = df.fillna(0)`

```
In [19]: df.isnull().any()
```

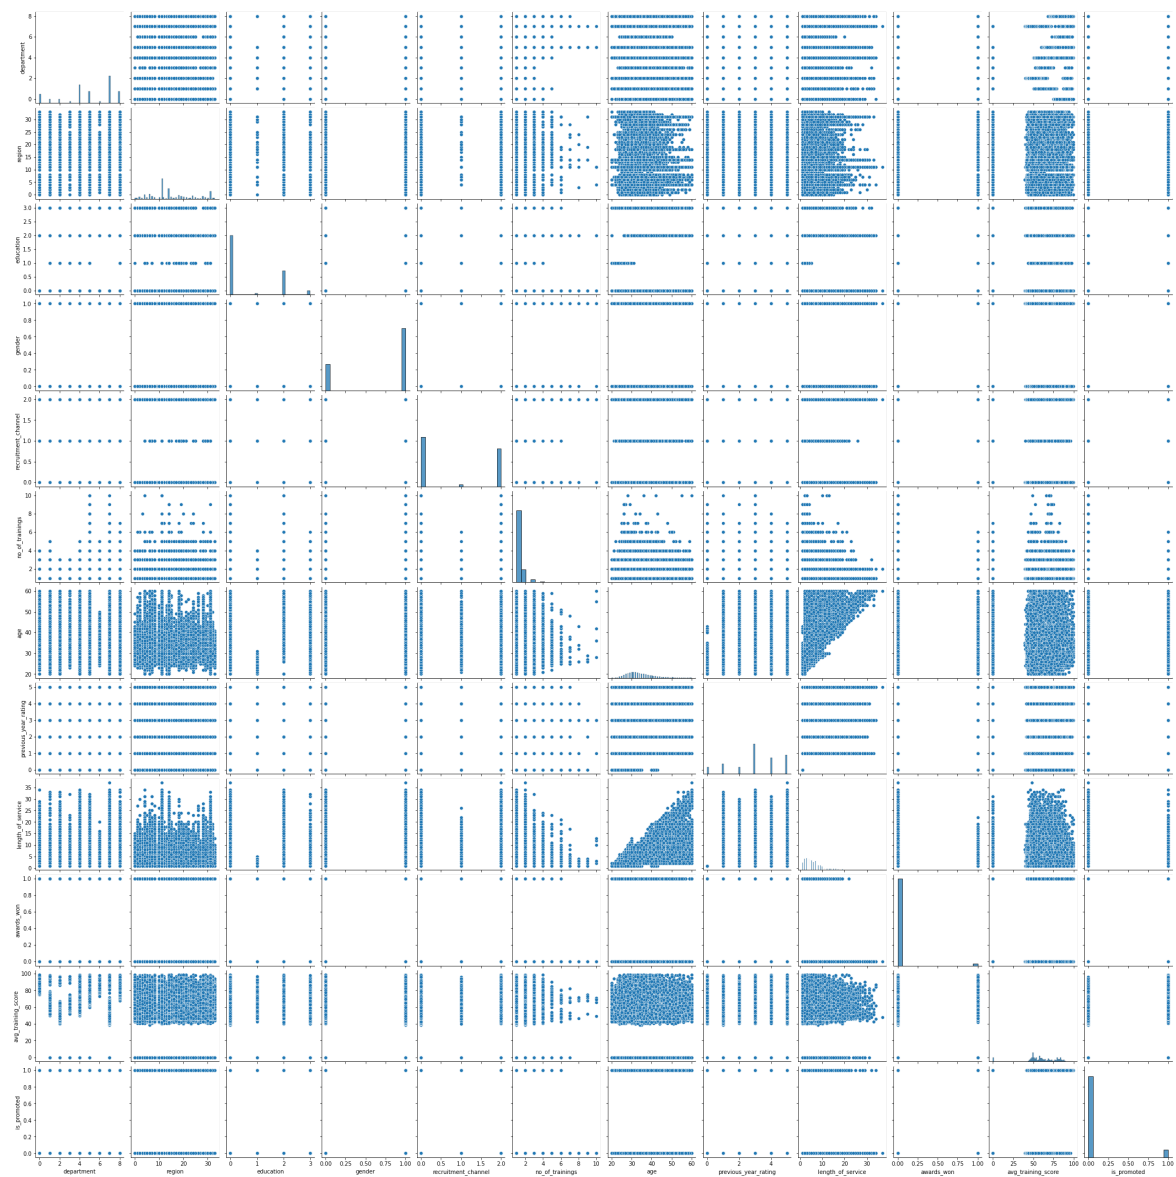
```
Out[19]: department      False
region                  False
education               False
gender                 False
recruitment_channel    False
no_of_trainings        False
age                   False
previous_year_rating   False
length_of_service      False
awards_won             False
avg_training_score     False
is_promoted            False
dtype: bool
```

```
In [20]: df['region'].value_counts()
```

```
Out[20]: 11      12343
14      6428
31      4843
6       2808
4       2648
18      2260
24      1935
28      1703
19      1659
7       1465
20      1318
2       1315
15      1175
21      994
25      945
10      874
12      850
5       827
17      819
8       796
29      766
30      690
23      657
32      655
1       648
0       610
16      508
3       500
33      420
13      411
22      346
27      292
26      269
9        31
Name: region, dtype: int64
```

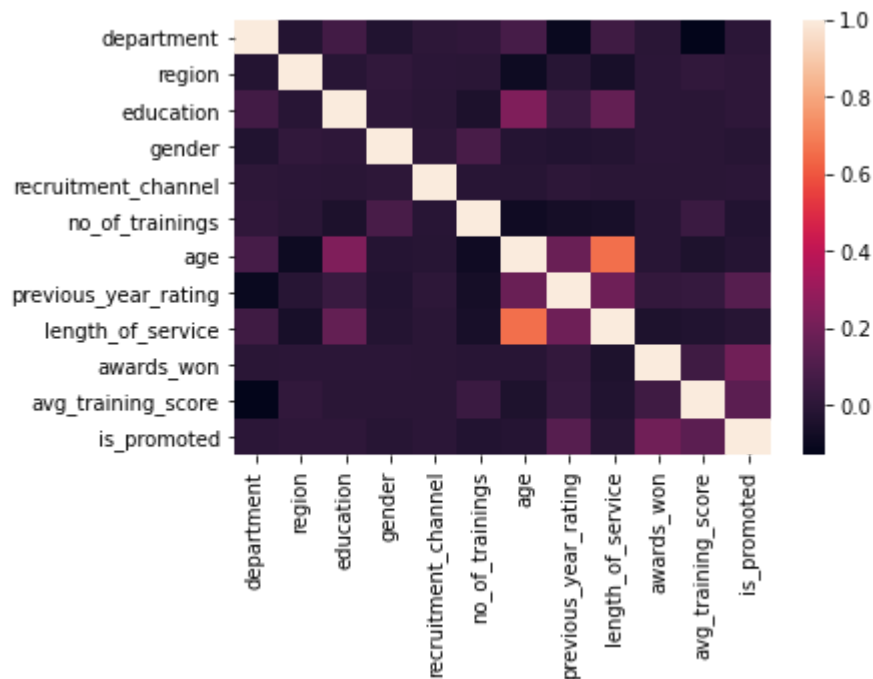
```
In [21]: sns.pairplot(df)
```

```
Out[21]: <seaborn.axisgrid.PairGrid at 0x23a5d9ce250>
```



```
In [22]: tc = df.corr()  
sns.heatmap(tc)
```

Out[22]: <AxesSubplot:>



```
In [23]: df.columns
```

Out[23]: Index(['department', 'region', 'education', 'gender', 'recruitment_channel',
'no_of_trainings', 'age', 'previous_year_rating', 'length_of_service',
'awards_won', 'avg_training_score', 'is_promoted'],
dtype='object')

```
In [24]: X = df[['department', 'region', 'education', 'gender', 'recruitment_channel',  
                'no_of_trainings', 'age', 'previous_year_rating', 'length_of_service',  
                'awards_won', 'avg_training_score']]  
Y = df[['is_promoted']]
```



```
In [25]: ▶ print(X.shape)
          print(Y.shape)
```

```
(54808, 11)
(54808, 1)
```

Splitting the dataset as train and test(test size of 30%)

```
In [26]: ▶ X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3,random_s
```

```
In [27]: ▶ print(X_train.shape)
          print(Y_train.shape)
```

```
(38365, 11)
(38365, 1)
```

Building the classification model

```
In [28]: ▶ from sklearn.linear_model import LogisticRegression
```

```
In [29]: ▶ model_lr = LogisticRegression()
```

```
In [30]: ▶ model_lr.fit(X_train,Y_train)
```

```
C:\xamp\anaconda1\lib\site-packages\sklearn\utils\validation.py:63: DataCon
versionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples, ), for example using ravel().
    return f(*args, **kwargs)
C:\xamp\anaconda1\lib\site-packages\sklearn\linear_model\_logistic.py:763:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[30]: LogisticRegression()
```

```
In [31]: ▶ pred = model_lr.predict(X_test)
```

```
In [32]:  pred
```

```
Out[32]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [33]:  from sklearn import metrics
```

```
In [34]:  print('The accuracy is ',(metrics.accuracy_score(Y_test,pred)*100),'%')
```

```
The accuracy is  91.60737091771574 %
```

```
In [35]:  from sklearn.tree import DecisionTreeClassifier
```

```
In [36]:  model = DecisionTreeClassifier()
```

```
In [37]:  model.fit(X_train,Y_train)
```

```
Out[37]: DecisionTreeClassifier()
```

```
In [39]:  Pred = model.predict(X_test)
```

```
In [40]:  print('The accuracy is ',(metrics.accuracy_score(Y_test,Pred)*100),'%')
```

```
The accuracy is  87.71513714042449 %
```

```
In [41]:  pip install xgboost
```

```
Collecting xgboostNote: you may need to restart the kernel to use updated p
ackages.
```

```
Using cached xgboost-1.4.2-py3-none-win_amd64.whl (97.8 MB)
Requirement already satisfied: scipy in c:\xamp\anaconda1\lib\site-packages
(from xgboost) (1.6.2)
Requirement already satisfied: numpy in c:\xamp\anaconda1\lib\site-packages
(from xgboost) (1.20.1)
Installing collected packages: xgboost
Successfully installed xgboost-1.4.2
```

```
In [42]:  import xgboost
```

```
In [43]:  from xgboost import XGBClassifier
```

```
In [44]:  model_xg = XGBClassifier()
```

```
In [45]: ► model_xg.fit(X_train,Y_train)
```

```
C:\xamp\anaconda1\lib\site-packages\xgboost\sklearn.py:1146: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].
```

```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

```
C:\xamp\anaconda1\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
return f(*args, **kwargs)
```

```
[11:33:02] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
```

```
Out[45]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                      colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                      importance_type='gain', interaction_constraints='',
                      learning_rate=0.300000012, max_delta_step=0, max_depth=6,
                      min_child_weight=1, missing=nan, monotone_constraints='()',
                      n_estimators=100, n_jobs=4, num_parallel_tree=1, random_state=0,
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                      tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [46]: ► predicted = model_xg.predict(X_test)
```

```
In [47]: ► print('The accuracy is ',(metrics.accuracy_score(Y_test,predicted)*100), '%')
```

```
The accuracy is 94.16164933406313 %
```

```
In [ ]: ►
```